

Question Answering using Structured and Semi-structured User Generated Content

Denis Savenkov
Emory University
Atlanta, USA
denis.savenkov@emory.edu

ABSTRACT

Modern question answering (QA) systems employ a variety of different unstructured (text corpora), semi-structured (question answer pairs, infoboxes) and structured (knowledge bases - KB) data sources for candidate generation. However, typically each of the data sources is processed independently and the merging of the results happens only after the candidates are generated. I propose to enrich the expressiveness of data representation of each of the data sources with complementary information for joint reasoning. Text documents can be annotated with the information about properties and connections between mentioned entities. On the other hand for knowledge base question answering systems passages and text fragments retrieved for a candidate answer can help the system to find a good mapping between a natural language question and the corresponding structured query representation.

Categories and Subject Descriptors

H.3.3 [H.3.3 Information Search and Retrieval]: Retrieval models; I.2.3 [Artificial Intelligence]: Answer/reason extraction

Keywords

Question Answering; Knowledge Bases; Web Search

1. INTRODUCTION

The ability to answer user questions with precise and concise information is a hard problem with a long history of research. Various data sources are available for candidate answer generation, two major ones are unstructured text corpora, and structured knowledge bases (e.g. dpPedia [3] and Freebase [8]). A hybrid approach to question answering [5, 16] generates candidates from multiple sources, however each of them is typically processed separately and results are merged on the scoring and ranking stage when some information is already lost. Efficient combination of different

information sources has potential to improve both text and knowledge base question answering systems. I propose to combine all the available sources together and do joint reasoning to generate better answer candidates and improve the overall question answering performance.

Question answering from text corpora typically starts by retrieving a set of potentially relevant documents using the question (or some transformation of the question [1]) as the query, and then extracting entities, phrases, sentences or paragraphs believed to be the answer to the question. However, the information available in the retrieved pieces of text is very limited and often not enough to decide whether it can be the answer to the given question. For example, below is one of the questions from TREC QA 2007 dataset:

“What republican senators supported the nomination of Harriet Miers to the Supreme Court?”

A candidate answer sentence *“Minority Leader Harry Reid had already offered his open support for Miers.”* mentions a senator “Harry Reid” and clearly says about his support of the nomination. However, “Harry Reid” is not a correct answer to the question because he is a member of the Democratic party. This information is not available in the answer candidate sentence, but it is present as one of the properties in Freebase: [Harry Reid, political_party, Democratic party]¹. Therefore, by looking into the knowledge available about the mentioned entities a QA system can make a better judgment about the candidate answer.

Question answering over linked data (knowledge bases) converts a natural language question into a structured query, such as SPARQL. The main challenge for such systems is to map words and phrases from the question to the corresponding entities and predicates from a KB. Usually, such lexicon is built during training using ground truth question-query pairs [12] or question-answer pairs [6]. Improvements were made by extending the lexicon using Wikipedia and patterns expressing certain predicates obtained via distant supervision [4, 9, 25, 28, 31]. But still, the amount of available labeled or weakly labeled training data is much smaller than the amount of unstructured data. This unstructured data will complement the learned lexicon, e.g. even if a question about a certain predicate wasn't seen during training, a set of text paragraphs mentioning both of the related entities can provide a QA system with enough evidence to make the correct decision.

Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. To copy otherwise, to republish, to post on servers or to redistribute to lists, requires prior specific permission and/or a fee.

WSDM '2016 San Francisco, USA

Copyright 20XX ACM X-XXXXX-XX-X/XX/XX ...\$15.00.

¹Actually, in Freebase the entities are connected by a path of length 2 through a mediator node. The predicates on the path are: /government/politician/party and /government/political_party_tenure/party

2. RELATED WORK

The field of question answering has a long history of research and dates back to 60s (see [19] for a survey of different approaches). The modern era of question answering research started with the rise of the Internet and exponential growth of information available in the World Wide Web. Since 1999 the annual TREC organized a number of open domain question answering shared tasks [13]. Approaches proposed over the years can be largely classified by the type of the information used to find the answers into knowledge base and text-based systems.

2.1 Text-based QA

A traditional approach to factoid question answering over document collections popularized by the TREC QA track is to retrieve a set of potentially relevant documents, extract and rank mentioned entities as candidate answers. One of the main challenges of such an approach is limited amount of information present in the extracted pieces of text. Systems test answer for incorrectness by matching the expected answer type with the type of candidate entity often predicted by an named entity tagger. These systems rely heavily on special complicated ontologies that encode the relationships between different question and answer types, e.g. [18, 20, 24]. Alternatively, the AskMSR system [10] (recently reviewed in [29]) used the redundancy of large text collections such as the web to extract n-grams that occur frequently in a retrieved set of documents. Their counting-based approach performed unexpectedly well on TREC 2001 and sparked an interest in exploring the web for question answering purposes [21]. However, in many cases the information from the extracted text fragments is not enough to make a judgment on an answer candidate. To solve this problem researchers experimented with using external resources, both unstructured (e.g. Wikipedia articles [2, 11]) and structured (e.g. Wordnet [23]), and demonstrated improved question answering performance. Recently [27] proposed to link entities from candidate answers to Freebase and use its type system and textual entity description for candidate scoring. However, most of the information in a KB is stored as relations between entities, therefore there is a big potential in using all available KB data to improve question answering.

2.2 Knowledge base QA

Recent development of large scale knowledge bases (e.g. dbPedia [3]) and Freebase [8]) motivated research in open domain question answering over linked data. Developed models can be compared on the annual QALD shared task ² and on a number of available benchmark datasets, e.g. WebQuestions [6]. The main challenge of such systems is to map natural language questions into the structured query representation. Such a lexicon can be learned from a labeled training set [6], ClueWeb collection aligned to Freebase [25, 31], question paraphrases clusters from WikiAnswers [7], Freebase triples rephrased as questions [9], and can be based on the embeddings of questions and knowledge base entities and predicates [9, 28]. However, most of the models are still biased towards the types of questions present in the training set and would benefit from more training data. In this work I propose to extend the training set with question-answer

²<http://greententacle.techfak.uni-bielefeld.de/~cunger/qald/>

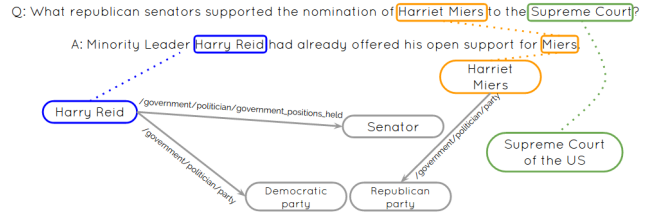


Figure 1: Annotation of natural language text with mentioned entities and their subgraphs in a knowledge base

pairs available on CQA websites, which were shown to be useful for relation extraction [26]. In addition, I propose to use unlabeled text resources for candidate query ranking, which can help to generalize to unseen types of questions and questions about predicates never mentioned in the training set.

2.3 Hybrid techniques

Hybrid question answering systems combine multiple available information sources, in particular text document and knowledge bases. Examples of such systems include IBM Watson [16], OpenQA [15], YodaQA [5]. The main difference between such systems and the proposed research is that hybrid systems typically use separate pipelines to extract candidates from different sources and only merge the candidate set while ranking. I propose to extend the representation of each of the data sources for better candidate generation from the beginning.

3. RESEARCH PROPOSAL

In this work I propose to enrich the input data representation for QA systems by combining available unstructured, semi-structured and structured data sources for joint reasoning, which can improve the performance of question answering over both text collections and knowledge bases.

3.1 Text-based QA

For question answering over text corpora I propose to extend the text representation with annotations about mentioned entities and their relations from open [15] or schema-based knowledge bases (e.g. dbPedia or Freebase). Such representation allows not only find different mentions of the same entity, but also look into the connections of the mentioned entities in order to learn more about the candidate answer. For example, for the question mentioned in the introduction "What republican senators supported the nomination of Harriet Miers to the Supreme Court?" and a candidate answer sentence "Minority Leader Harry Reid had already offered his open support for Miers.", such joint text-KB representation can look like Figure 1. A QA system can discover that "Harry Reid" political affiliation is with the Democratic Party, and he cannot be referred to as "republican senator". In other cases using a KB as an additional source of information may reveal specific connections between entities in the question and in the answer candidates. For example, for another TREC QA 2007 question "For which newspaper does Krugman write?" and retrieved candidate answer *New York Times* a path between "Paul Krugman" and "New York Times" in the knowledge graph gives an evidence in support

Table 1: Example of a question from WebQuestions dataset with related unstructured information

Question	Who is the woman that John Edwards had an affair with?
Provided answer	“Writer”, “Politician”, “Lawyer”, “Attorneys in the United States”
Correct answer	Rielle Hunter
Phrase from Wikipedia	John Edwards had engaged in an affair with Rielle Hunter...
QnA pair from Yahoo! Answers	Who was it that John Edwards had an affair with? Today, John Edwards admitted to having an affair with filmmaker Rielle Hunter.

of the candidate.

More specifically, to do this kind of inference I propose:

- use existing approaches for document retrieval (e.g. web search using question as a query [29]) and candidate answer extraction.
- perform entity linking to mentions of KB entities in questions and corresponding candidate answers.
- for each mentioned entity extract a subgraph containing its neighborhood up to certain number of edges away and paths to other mentioned entities.
- follow machine learning approach for candidate answer ranking and extend the feature representation with features derived from subgraph analysis. Examples of features are:
 - features describing discovered connections between entities mentioned in a question and a candidate answer, such as indicators of the relations, combination of relations with words and n-grams from the questions, similarity between the relations and the question text (using tf-idf or embeddings representation), etc. Textual representations of the predicates in structured knowledge bases can be obtained either from its description or using patterns learned from a large collection using distant supervision [22].
 - features describing the entities mentioned in the answer, i.e. similarities between entity properties and question words, n-grams and phrases, etc.

For training text-based QA model I propose to use available QnA pairs from community question answering websites, which represent real user tasks and after certain filtering can be a good fit for training both factoid and non-factoid question answering systems. The data can help to learn more associations between the language used in questions and their corresponding answers, which can be encoded as conditional probabilities (e.g. $p(w_a | w_{q_1}, \dots, w_{q_n})$, where w_a is a word of the answer and w_{q_i} is some subset of the question words), pointwise mutual information or by employing deep learning techniques [30].

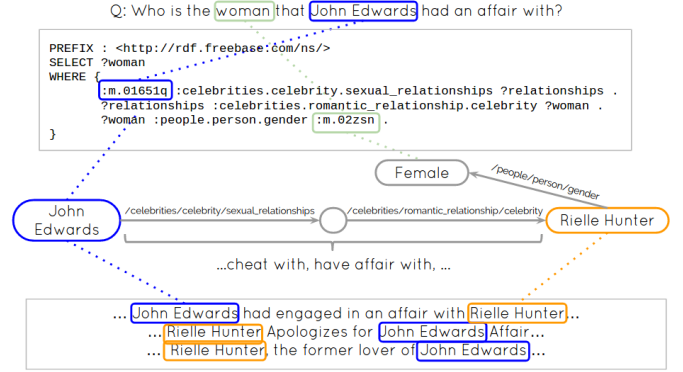


Figure 2: Annotation of KB graph nodes and edges with unstructured text data

3.2 Knowledge base QA

Lexicons learned during training of a knowledge base question answering systems are limited and often needs to be retrained to include additional data. To complement the lexicon learned during training for candidate structured query scoring I propose to use unstructured text data related to mentioned entities and predicates (see Figure 2). For example, Table 1 shows an example of a question from the WebQuestions dataset, that is answered incorrectly by a state-of-the-art system. A similar question is missing from the training set, however, an easy web search can retrieve a relevant sentence, which give enough supporting evidence to answer this question correctly.

More specifically, I propose:

- Use one of the available state of the art systems, such as [4], as a baseline.
- Extend a set of features representing a candidate answer with features derived from unstructured text fragments:
 - from large document collection (such as the web) retrieve a set of passages by querying a search system with question, question + answer, question and answer entities as queries.
 - find mentions of answer entities in the passages and use some aggregated statistics as features for the corresponding candidates.
 - for each candidate answer retrieve a set of patterns used to express the corresponding predicates obtained using distant supervision from a large text collection and compute the similarities between these patterns and the question text.

In addition, similar to text-based systems, I propose to include QnA pairs from CQA websites as weakly labeled training data. For example, a question similar to the one presented in Table 1 is not present in the labeled training set, but can easily be found in Yahoo! Answers³. CQA data should be preprocessed (possibly filtered by categories of the questions and other heuristics), and select QnA pairs mentioning at least one entity in the question and in the answer. After a reasonable cleanup such QnA pairs can

³<http://answers.yahoo.com/>

be used for training treating answer entities as the correct answer.

4. EXPERIMENTS

4.1 Text-based QA

TREC QA datasets served as a benchmark for various question answering systems. Therefore, to evaluate the proposed approach for question answering over text enriched with the structured data I propose to test it on dataset derived from TREC QA and compare against existing strong baselines, including the most related approaches [15, 27]. The proposed system can use the web as the corpus and query it using Bing Search API⁴. Freebase and Reverb extractions [14] are examples of schema-based and open knowledge bases that can be used for the experiments. The metrics used for evaluation typically include accuracy and mean reciprocal rank (MRR).

For non-factoid question answering this year TREC pioneered a new question answering track - TREC LiveQA⁵, which targets questions asked by real users of Yahoo! Answers website. This year the deadline for system submission was on August 31 and my system trained on CQA QnA pairs participated in the challenge. The results will be available on the TREC Conference in November 2015. Organizers plan to continue with another TREC LiveQA task next year and this is going to be a good estimation of the effectiveness of the proposed techniques on hard real user questions.

4.2 Knowledge base QA

Most of the recent work on knowledge base question answering and semantic parsing have been evaluated on the WebQuestions dataset [6], which contains a collection of question text and correct answer entities. The questions were collected using Google Suggest API and answers crowdsourced using Amazon Mechanical Turk⁶. The proposed approach will be compared against the previous results⁷ on this dataset. Again, web can be used as a text collection which can be queried using Bing Search API. Relation extraction patterns can be mined using distant supervision from ClueWeb collection using publicly available dataset of Freebase annotations [17].

However, WebQuestions dataset has certain limitations, e.g. questions mined using Google Suggest API have very similar structure and lexicon, and to find the answer to the mined questions users were asked to use the question entity Freebase profile page, which only include entities connected directly with a predicate or through a mediator node. Therefore most of the state-of-the-art results on the dataset use a small number of predefined logical form patterns. On the other hand CQA websites have a fraction of factoid questions with provided text answers. Here I propose to use to construct a new dataset for question answering over Freebase by selecting a subset of QnA pairs with at least one entity in question and answer and some reasonable filtering heuristics and manual validation using crowdsourcing (e.g. through Amazon Mechanical Turk). Existing systems need

to be retrained and tested on the new dataset to compare against the proposed model.

5. DISCUSSION

The proposed approach targets the problem of improving the performance of question answering systems using joint reasoning over unstructured, semi-structured and structured data sources. By linking entity mentions to their knowledge base objects a text-based QA system will be able to use not only lexical information present in extracted text fragments, but also all the factual information about the entities, which should improve its performance. On the other hand, knowledge base question answering should benefit from textual data about predicates and entities mentioned in a questions and a candidate answer. Additional unstructured data will serve as a bridge between a natural language question and the corresponding knowledge base query, which should boost the recall of question answering systems.

However, there are certain questions and limitations, that I would like to discuss. As we know, knowledge bases are inherently incomplete: not only many facts are missing, but also a set of predicates is far from being complete. Therefore, for many questions there are no corresponding predicates in a knowledge base. Given the fact that at the moment text-based QA systems outperform knowledge base systems on factoid questions from the TREC QA dataset, it is unclear how much additional information a KB can add and how big is an advantage over hybrid approaches that simply combine the candidates obtained from various data sources. An alternative approach to get more knowledge about candidate answers is to retrieve more unstructured data, e.g. previous research found Wikipedia articles to be useful. Another question is related to the usefulness of the information stored in a KB for complex and non-factoid questions. The main challenge is to “understand” the text of the answer and predict whether it replies to the question. Facts stored in Freebase or similar KB might not reveal much about the meaning of the answer and we would need a different source of knowledge.

6. REFERENCES

- [1] E. Agichtein, S. Lawrence, and L. Gravano. Learning search engine specific query transformations for question answering. In *Proceedings of the Tenth International World Wide Web Conference, WWW 10*, pages 169–178, 2001.
- [2] D. Ahn, V. Jijkoun, G. Mishne, K. Müller, M. de Rijke, and K. Schlobach. Using wikipedia at the trec qa track. In *The Thirteenth Text Retrieval Conference (TREC 2004)*, 2005.
- [3] S. Auer, C. Bizer, G. Kobilarov, J. Lehmann, R. Cyganiak, and Z. Ives. *Dbpedia: A nucleus for a web of open data*. Springer, 2007.
- [4] H. Bast and E. Haussmann. More accurate question answering on freebase. In *CIKM*, 2015.
- [5] P. Baudiš and J. Šedivý. Modeling of the question answering task in the yodaqa system. In *Experimental IR Meets Multilinguality, Multimodality, and Interaction*, pages 222–228. Springer, 2015.
- [6] J. Berant, A. Chou, R. Frostig, and P. Liang. Semantic parsing on freebase from question-answer pairs. In *Proceedings of the 2013 Conference on*

⁴<https://datamarket.azure.com/dataset/bing/searchweb>

⁵<http://trec-liveqa.org/>

⁶<http://mturk.com/>

⁷<http://goo.gl/sePBja>

- Empirical Methods in Natural Language Processing*, pages 1533–1544, 2013.
- [7] J. Berant and P. Liang. Semantic parsing via paraphrasing. In *Proceedings of the 52nd Annual Meeting of the Association for Computational Linguistics, ACL 2014*, pages 1415–1425, 2014.
 - [8] K. Bollacker, C. Evans, P. Paritosh, T. Sturge, and J. Taylor. Freebase: A collaboratively created graph database for structuring human knowledge. In *Proceedings of the 2008 ACM International Conference on Management of Data, SIGMOD '08*, pages 1247–1250, New York, NY, USA, 2008. ACM.
 - [9] A. Bordes, S. Chopra, and J. Weston. Question answering with subgraph embeddings. In *Proceedings of the 2014 Conference on Empirical Methods in Natural Language Processing, EMNLP 2014*, pages 615–620, 2014.
 - [10] E. Brill, J. Lin, M. Banko, S. Dumais, and A. Ng. Data-intensive question answering. In *Proceedings of TREC 2001*, January 2001.
 - [11] D. Buscaldi and P. Rosso. Mining knowledge from wikipedia for the question answering task. In *Proceedings of the International Conference on Language Resources and Evaluation*, pages 727–730, 2006.
 - [12] Q. Cai and A. Yates. Large-scale semantic parsing via schema matching and lexicon extension. In *Proceedings of the 51st Annual Meeting of the Association for Computational Linguistics, ACL 2013*, pages 423–433, 2013.
 - [13] H. T. Dang, D. Kelly, and J. J. Lin. Overview of the trec 2007 question answering track. In *TREC*, volume 7, page 63. Citeseer, 2007.
 - [14] A. Fader, S. Soderland, and O. Etzioni. Identifying relations for open information extraction. In *Proceedings of the 2011 Conference on Empirical Methods in Natural Language Processing, EMNLP 2011*, pages 1535–1545, 2011.
 - [15] A. Fader, L. Zettlemoyer, and O. Etzioni. Open question answering over curated and extracted knowledge bases. In *Proceedings of the 20th ACM International Conference on Knowledge Discovery and Data Mining, KDD '14*, pages 1156–1165, New York, NY, USA, 2014. ACM.
 - [16] D. A. Ferrucci, E. W. Brown, J. Chu-Carroll, J. Fan, D. Gondek, A. Kalyanpur, A. Lally, J. W. Murdock, E. Nyberg, J. M. Prager, N. Schlaefer, and C. A. Welty. Building watson: An overview of the deepqa project. *AI Magazine*, 31(3):59–79, 2010.
 - [17] E. Gabrilovich, M. Ringgaard, and A. Subramanya. Facc1: Freebase annotation of cluweb corpora, 2013.
 - [18] E. H. Hovy, L. Gerber, U. Hermjakob, M. Junk, and C.-Y. Lin. Question answering in webclopedia. In *TREC*, volume 52, pages 53–56, 2000.
 - [19] O. Kolomiyets and M.-F. Moens. A survey on question answering technology from an information retrieval perspective. *Information Sciences*, 181(24):5412–5434, Dec. 2011.
 - [20] X. Li and D. Roth. Learning question classifiers. In *19th International Conference on Computational Linguistics, COLING 2002*, 2002.
 - [21] J. J. Lin and B. Katz. Question answering from the web using knowledge annotation and knowledge mining techniques. In *Proceedings of the 2003 ACM CIKM International Conference on Information and Knowledge Management*, pages 116–123, 2003.
 - [22] M. Mintz, S. Bills, R. Snow, and D. Jurafsky. Distant supervision for relation extraction without labeled data. In *ACL 2009, Proceedings of the 47th Annual Meeting of the Association for Computational Linguistics and the 4th International Joint Conference on Natural Language Processing of the AFNLP*, pages 1003–1011, 2009.
 - [23] M. Pasca and S. Harabagiu. The informative role of wordnet in open-domain question answering. In *Proceedings of NAACL-01 Workshop on WordNet and Other Lexical Resources*, pages 138–143, 2001.
 - [24] J. Prager, J. Chu-Carroll, E. W. Brown, and K. Czuba. Question answering by predictive annotation. In *Advances in Open Domain Question Answering*, pages 307–347. Springer, 2006.
 - [25] S. Reddy, M. Lapata, and M. Steedman. Large-scale semantic parsing without question-answer pairs. *TACL*, 2:377–392, 2014.
 - [26] D. Savenkov, W. Lu, J. Dalton, and E. Agichtein. Relation extraction from community generated question-answer pairs. In *NAACL HLT 2015, The 2015 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies*, pages 96–102, 2015.
 - [27] H. Sun, H. Ma, W.-t. Yih, C.-T. Tsai, J. Liu, and M.-W. Chang. Open domain question answering via semantic enrichment. In *Proceedings of the 24th International Conference on World Wide Web, WWW '15*, pages 1045–1055, Republic and Canton of Geneva, Switzerland, 2015. International World Wide Web Conferences Steering Committee.
 - [28] W. tau Yih, M.-W. Chang, X. He, and J. Gao. Semantic parsing via staged query graph generation: Question answering with knowledge base. In *Proceedings of the Joint Conference of the 53rd Annual Meeting of the ACL and the 7th International Joint Conference on Natural Language Processing of the AFNLP. ACL & Association for Computational Linguistics*, July 2015.
 - [29] C. Tsai, W.-t. Yih, and C. Burges. Web-based question answering: Revisiting askmsr. Technical report, Technical Report MSR-TR-2015-20, Microsoft Research, 2015.
 - [30] D. Wang and E. Nyberg. A long short-term memory model for answer sentence selection in question answering. In *Proceedings of the 53rd Annual Meeting of the Association for Computational Linguistics and the 7th International Joint Conference on Natural Language Processing of the Asian Federation of Natural Language Processing, ACL 2015*, pages 707–712, 2015.
 - [31] X. Yao and B. V. Durme. Information extraction over structured data: Question answering with freebase. In *Proceedings of the 52nd Annual Meeting of the Association for Computational Linguistics, ACL 2014*, pages 956–966, 2014.

APPENDIX

A. STUDENT STATEMENT

I'm a 5th year PhD student at Emory University and during the course of my studies I've worked on a number of different topics, starting from click log analysis, study of mobile users behavior, search user assistance for difficult search tasks, information extraction and now question answering. In addition I've spent 4 wonderful summers in Microsoft and Google working on search results diversity, entity linking, information extraction and coreference resolution. From a wide range of topics I've worked on, I've decided to focus my thesis on question answering. In several month I'm going to do a thesis proposal and WSDM Doctoral Consortium presents a great chance to discuss the proposed direction with experienced researchers with various background. I hope to make a valuable contribution to the area and discussion with the experts will help me polish the plan for the future thesis. The proposed topic involves several related areas, such as information retrieval, natural language processing, web search and mining, knowledge representation, knowledge base construction and semantic parsing. Most of my work was in the information retrieval area, and I have less experience with knowledge representation and natural language processing. Therefore, I would greatly appreciate a chance to attend the consortium and discuss the potential challenges and ways to improve the work with researches with complementary experiences, which can make my work more interesting for different communities.

B. ADVISOR STATEMENT

My name is Eugene Agichtein⁸, and I am Denis Savenkov's Ph.D. advisor. Denis is entering his 5th year as a PhD student at the Emory Mathematics & Computer Science department. He is expected to present his thesis proposal in the next few months, and to defend his thesis within a year, likely by December 2016. Denis has broad interests within Information Retrieval, and has done excellent work on a number of topics in IR. In particular, he and his team placed in the top 3 teams at the Yandex Click prediction challenge, and subsequently won the "Search Engine Prediction" challenge, publishing a full SIGIR paper on the novel techniques they developed. Denis has also done in-depth studies of community question answering data, resulting in a full ICWSM paper as well as interesting analysis of how (not) to support searchers during complex tasks. The analysis of difficult, or failed, searches, led Denis to become interested in question answering, as a way of attacking some of the long standing challenges in search. I believe Denis has the skills and the creativity to tackle the ambitious task he proposes, on combining information from unstructured text and knowledge bases for question answering. Attending the doctoral consortium at WSDM will give Denis a chance to interact with leading experts in Web search, social media analysis and data mining, and especially those with complementary expertise to mine. As an IR researcher, I am reasonably confident on advising Denis on the aspects of his thesis related to IR and QA over unstructured (text) data, but am less familiar with the work done in KD and linked data communities, both in academia or industry. As Denis' thesis spans both areas, input from researchers who work on constructing

and utilizing KBs, knowledge graphs and other structured sources is critical. Participating in DC would be particularly beneficial to Denis this year, as he is at the stage of developing his thesis proposal.

⁸<http://www.mathcs.emory.edu/~eugene/>